


FINAL CAPSTONE PROJECT

CIFAR10 Dataset

Image Recognition using Convolutional Neural networks

Presented by: Usman Shaikh





Dataset Description

- Dataset: <https://www.cs.toronto.edu/~kriz/cifar.html>
- CIFAR-10 is a subset of the 80 million tiny images dataset created by **Canadian Institute for Advanced Research (CIFAR)**.
- The CIFAR-10 dataset consists of 60000 32x32x3 color images in 10 equal classes, (6000 images per class).
- Each class of images corresponds to a physical object (automobile, cat, dog, airplane, etc.)
- In total there are 10 classes including airplane, automobile, bird, cat, deer, dog, frog, horse, ship and a truck.

Problem Summary


- The purpose of this project is to develop a multi class classification model successfully predicting different classes.
- Our goal is to train and test different machine learning models as per our requirement and then come up with one specific model which provides us with the best possible and consistent results in a timely fashion.
- For this purpose we will be building eleven different machine learning models, two each for the Random Forest, Logistic Regression and Gradient Boost followed by 6 separate models of Convolutional Neural Networks.

RANDOM FOREST



Without PCA

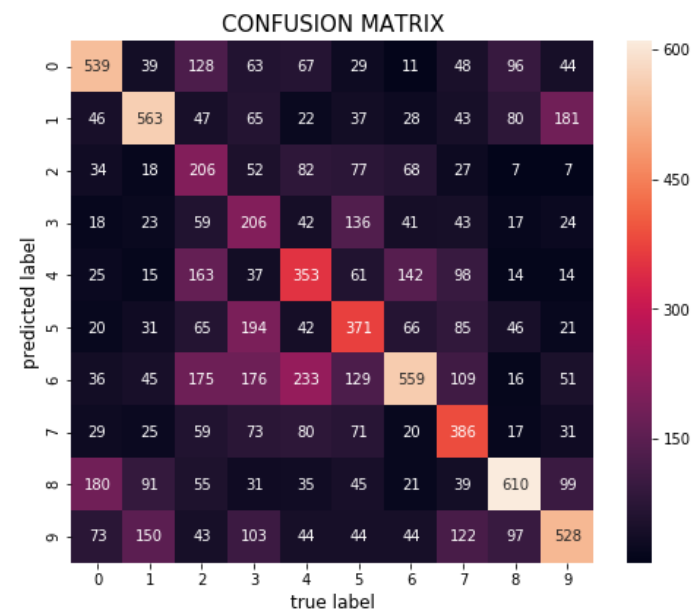
- Variables: 3,072
- Estimators: 500
- Max depth: 10
- R^2 Train Data: 0.64
- R^2 Test Data: 0.43
- Mean Absolute Error: 2.03
- Mean Squared Error: 10.67
- Root Mean Squared Error: 3.26
- Time: 00:08:16.12



Classification	Report		Testing Set		support
	precision	recall	f1-score		
0	0.52	0.52	0.52	1000	
1	0.49	0.50	0.50	1000	
2	0.41	0.19	0.26	1000	
3	0.34	0.18	0.23	1000	
4	0.33	0.43	0.38	1000	
5	0.40	0.38	0.39	1000	
6	0.39	0.57	0.46	1000	
7	0.45	0.42	0.44	1000	
8	0.54	0.58	0.56	1000	
9	0.43	0.56	0.49	1000	
micro avg	0.43	0.43	0.43	10000	
macro avg	0.43	0.43	0.42	10000	
weighted avg	0.43	0.43	0.42	10000	

With PCA

- Components = 56
- Estimators: 500
- Max depth: 10
- R^2 Train Data: 0.60
- R^2 Test Data: 0.43
- Mean Absolute Error: 2.05
- Mean Squared Error: 10.66
- Root Mean Squared Error: 3.26
- Time: 00:01:53.35




Classification	Report precision	Testing Set recall	f1-score	support
0	0.51	0.54	0.52	1000
1	0.51	0.56	0.53	1000
2	0.36	0.21	0.26	1000
3	0.34	0.21	0.26	1000
4	0.38	0.35	0.37	1000
5	0.39	0.37	0.38	1000
6	0.37	0.56	0.44	1000
7	0.49	0.39	0.43	1000
8	0.51	0.61	0.55	1000
9	0.42	0.53	0.47	1000
micro avg	0.43	0.43	0.43	10000
macro avg	0.43	0.43	0.42	10000
weighted avg	0.43	0.43	0.42	10000

LOGISTIC REGRESSION



Without PCA


- Variables: 3,072
- R^2 Train Data: 0.52
- R^2 Test Data: 0.36
- Mean Absolute Error: 2.31
- Mean Squared Error: 12.20
- Root Mean Squared Error: 3.49
- Time: 01:49:33.14



Classification	Report	Testing Set			
	precision	recall	f1-score	support	
0	0.42	0.45	0.43	1000	
1	0.41	0.44	0.42	1000	
2	0.28	0.25	0.26	1000	
3	0.24	0.23	0.24	1000	
4	0.34	0.28	0.31	1000	
5	0.30	0.29	0.30	1000	
6	0.38	0.42	0.40	1000	
7	0.41	0.42	0.41	1000	
8	0.46	0.49	0.48	1000	
9	0.41	0.42	0.42	1000	
micro avg	0.37	0.37	0.37	10000	
macro avg	0.37	0.37	0.37	10000	
weighted avg	0.37	0.37	0.37	10000	

With PCA

- Components = 56
- R^2 Train Data: 0.3832
- R^2 Test Data: 0.38
- Mean Absolute Error: 2.25
- Mean Squared Error: 11.86
- Root Mean Squared Error: 3.44
- Time: 00:00:16.38



Classification	Report Testing Set				
		precision	recall	f1-score	support
	0	0.44	0.46	0.45	1000
	1	0.42	0.49	0.45	1000
	2	0.29	0.24	0.26	1000
	3	0.28	0.22	0.25	1000
	4	0.35	0.26	0.30	1000
	5	0.35	0.33	0.34	1000
	6	0.38	0.50	0.43	1000
	7	0.40	0.40	0.40	1000
	8	0.44	0.49	0.47	1000
	9	0.42	0.46	0.44	1000
	micro avg	0.38	0.38	0.38	10000
	macro avg	0.38	0.38	0.38	10000
	weighted avg	0.38	0.38	0.38	10000

GRADIENT BOOST





Without PCA

- Variables: 3,072
- R^2 Train Data: 0.94
- R^2 Test Data: 0.44
- Mean Absolute Error: 1.98
- Mean Squared Error: 10.21
- Root Mean Squared Error: 3.19
- Subsample = 0.3
- Estimators = 50
- Max depth = 10
- Time: 09:24:40.93

With PCA

- Components = 56
- R^2 Train Data: 0.92
- R^2 Test Data: 0.44
- Mean Absolute Error: 2.01
- Mean Squared Error: 10.35
- Root Mean Squared Error: 3.21
- Subsample = 0.3
- Estimators = 50
- Max depth = 10
- Time: 00:14:35.99



Classification Report Testing Set				
	precision	recall	f1-score	support
0	0.51	0.51	0.51	1000
1	0.52	0.52	0.52	1000
2	0.33	0.35	0.34	1000
3	0.27	0.28	0.27	1000
4	0.42	0.39	0.40	1000
5	0.38	0.35	0.36	1000
6	0.46	0.53	0.49	1000
7	0.46	0.43	0.44	1000
8	0.57	0.59	0.58	1000
9	0.48	0.47	0.48	1000
micro avg	0.44	0.44	0.44	10000
macro avg	0.44	0.44	0.44	10000
weighted avg	0.44	0.44	0.44	10000


MODEL COMPARISON



	Random Forest	Logistic Regression	Gradient boost
Without PCA	R ² Train Data: 0.64 R ² Test Data: 0.43	R ² Train Data: 0.52 R ² Test Data: 0.36	R ² Train Data: 0.94 R ² Test Data: 0.44
Time	00:08:16.12	01:49:33.14	09:24:40.93
With PCA	R ² Train Data: 0.60 R ² Test Data: 0.43	R ² Train Data: 0.38 R ² Test Data: 0.38	R ² Train Data: 0.92 R ² Test Data: 0.44
Time	00:01:53.35	00:00:16.38	00:09:43.77



- For comparison model parameters are kept same for both instances.
- It shows that we are better off using Principle Component Analysis (PCA).
- Processing times can be considerably reduced without compromising accuracy.
- Models still require further tuning and fail to produce desired results.



CONVOLUTIONAL NEURAL NETWORKS – MODEL 1




MODEL 1 (4 Layers)

Name: model1



- Variables: 3,072
- Convolutional layers: 4
- Max pooling layers: 2
- Dense Layers: 2
- Dropout layers: 3 (0.25 – 0.5)
- Activation Functions used
 - Rectified Linear Unit (RELU)
 - Soft max
- Optimizer: SGD
- Loss = sparse categorical cross entropy
- Learning rate = 0.0001
- Batch size = 32
- Epochs = 100
- Total time: 02:29:19.31

Test loss: 1.72

Test accuracy: 0.38



Classification Report					
	precision	recall	f1-score	support	
0	0.45	0.50	0.48	1000	
1	0.48	0.56	0.52	1000	
2	0.29	0.03	0.06	1000	
3	0.23	0.14	0.17	1000	
4	0.30	0.41	0.35	1000	
5	0.37	0.24	0.29	1000	
6	0.35	0.59	0.44	1000	
7	0.36	0.48	0.42	1000	
8	0.43	0.49	0.46	1000	
9	0.45	0.38	0.41	1000	
micro avg	0.38	0.38	0.38	10000	
macro avg	0.37	0.38	0.36	10000	
weighted avg	0.37	0.38	0.36	10000	



CONVOLUTIONAL NEURAL NETWORKS – MODEL 2




MODEL 2 (4 Layers)

Name: *model1b*



- Variables: 3,072
- Convolutional layers: 4
- Max pooling layers: 2
- Dense Layers: 2
- Dropout layers: 3 (0.25 – 0.5)
- Activation Functions used
 - Rectified Linear Unit (RELU)
 - Soft max
- Optimizer: SGD
- Loss = sparse categorical cross entropy
- Learning rate = 0.001 +
- Batch size = 64 +
- Epochs = 100
- Total time: 02:18:07.45

Test loss: 1.25

Test accuracy: 0.55



Classification Report					
	precision	recall	f1-score	support	
0	0.62	0.53	0.57	1000	
1	0.67	0.74	0.70	1000	
2	0.45	0.27	0.33	1000	
3	0.36	0.41	0.38	1000	
4	0.46	0.38	0.42	1000	
5	0.46	0.58	0.51	1000	
6	0.63	0.67	0.65	1000	
7	0.57	0.62	0.59	1000	
8	0.65	0.69	0.67	1000	
9	0.66	0.64	0.65	1000	
micro avg	0.55	0.55	0.55	10000	
macro avg	0.55	0.55	0.55	10000	
weighted avg	0.55	0.55	0.55	10000	



CONVOLUTIONAL NEURAL NETWORKS – MODEL 3



MODEL 3 (6 Layers)


Name: *model2b*

- Variables: 3,072
- Convolutional layers: 6 +
- Max pooling layers: 3 +
- Dense Layers: 2
- Dropout layers: 4 (0.25 – 0.5) +
- Activation Functions used
 - Rectified Linear Unit (RELU)
 - Soft max
- Optimizer: SGD
- Loss = sparse categorical cross entropy
- Learning rate = 0.001
- Batch size = 32
- Epochs = 100 +
- Total time: 04:27:32.18

Test loss: 0.89

Test accuracy: 0.68

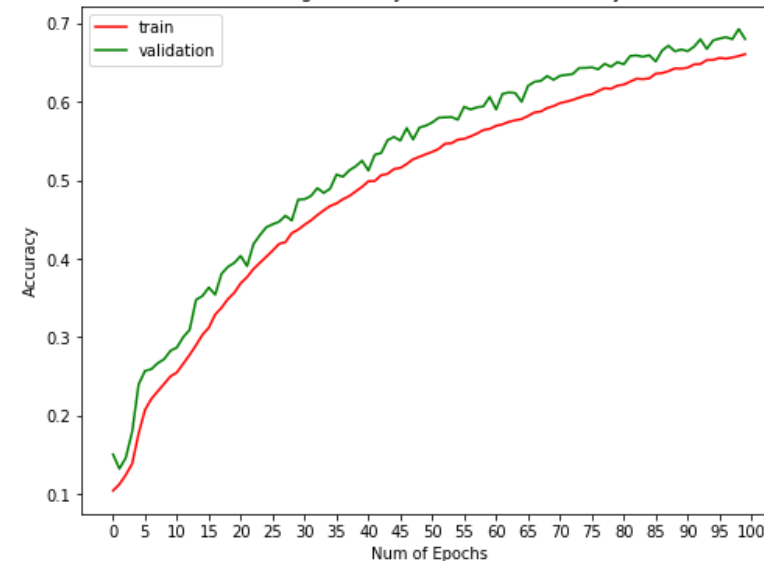
CONFUSION MATRIX




Classification Report

	precision	recall	f1-score	support
0	0.75	0.65	0.70	1000
1	0.82	0.85	0.83	1000
2	0.59	0.48	0.53	1000
3	0.53	0.40	0.45	1000
4	0.58	0.64	0.61	1000
5	0.61	0.55	0.58	1000
6	0.59	0.87	0.70	1000
7	0.76	0.72	0.74	1000
8	0.75	0.84	0.79	1000
9	0.80	0.80	0.80	1000
micro avg	0.68	0.68	0.68	10000
macro avg	0.68	0.68	0.67	10000
weighted avg	0.68	0.68	0.67	10000

Training Accuracy vs Validation Accuracy



Training Loss vs Validation Loss



CONVOLUTIONAL NEURAL NETWORKS – MODEL 4



MODEL 4 (8 Layers)


Name: *model3*

- Variables: 3,072
- Convolutional layers: 8 +
- Max pooling layers: 4
- Dense Layers: 2
- Dropout layers: 5 (0.25 – 0.5) +
- Activation Functions used
 - Rectified Linear Unit (RELU)
 - Soft max
- Optimizer: SGD
- Loss = sparse categorical cross entropy
- Learning rate = 0.0001 -
- Batch size = 64 +
- Epochs = 100
- Total time: 05:28:11.04

Test loss: 0.886

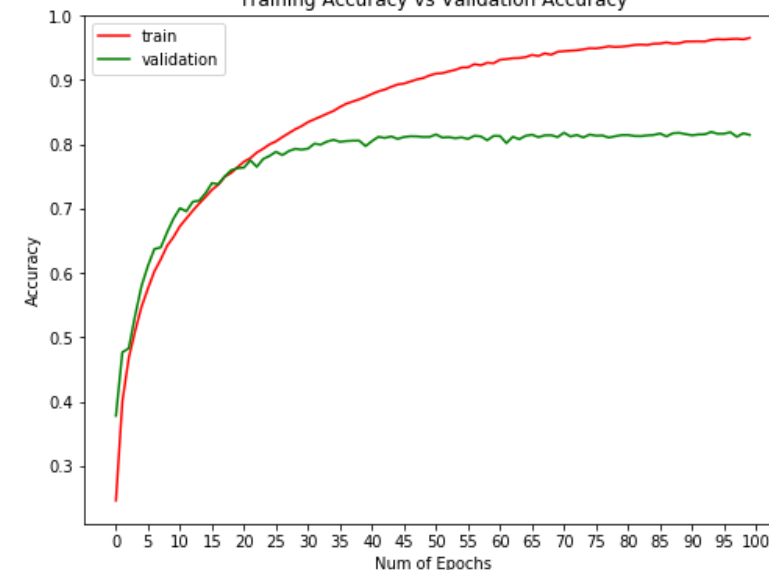
Test accuracy: 0.814

CONFUSION MATRIX




Classification Report				
	precision	recall	f1-score	support
0	0.81	0.87	0.84	1000
1	0.89	0.91	0.90	1000
2	0.79	0.72	0.75	1000
3	0.70	0.60	0.64	1000
4	0.80	0.78	0.79	1000
5	0.71	0.73	0.72	1000
6	0.86	0.85	0.86	1000
7	0.83	0.88	0.85	1000
8	0.89	0.89	0.89	1000
9	0.86	0.91	0.88	1000
micro avg	0.81	0.81	0.81	10000
macro avg	0.81	0.81	0.81	10000
weighted avg	0.81	0.81	0.81	10000

Training Accuracy vs Validation Accuracy



Training Loss vs Validation Loss



CONVOLUTIONAL NEURAL NETWORKS – MODEL 5



MODEL 5 (10 Layers)


Name: *model4*

- Variables: 3,072
- Convolutional layers: 10 +
- Max pooling layers: 4
- Dense Layers: 2
- Dropout layers: 5 (0.25 – 0.5) +
- Activation Functions used
 - Rectified Linear Unit (RELU)
 - Soft max
- Optimizer: SGD
- Loss = sparse categorical cross entropy
- Learning rate = 0.0001
- Batch size = 64
- Epochs = 100
- Total time: 10:30:20.44

Test loss: 0.830

Test accuracy: 0.816

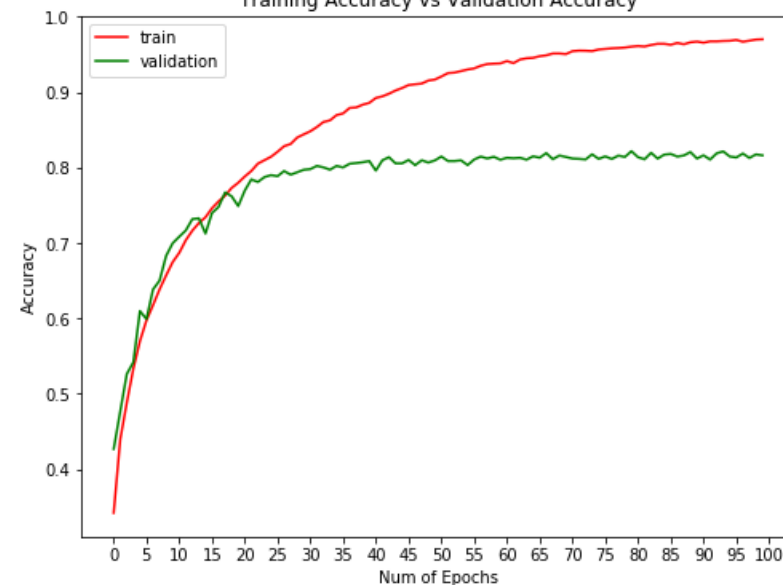
CONFUSION MATRIX




Classification Report

	precision	recall	f1-score	support
0	0.85	0.84	0.85	1000
1	0.92	0.92	0.92	1000
2	0.80	0.70	0.75	1000
3	0.63	0.67	0.65	1000
4	0.79	0.79	0.79	1000
5	0.72	0.68	0.70	1000
6	0.84	0.87	0.85	1000
7	0.84	0.87	0.85	1000
8	0.89	0.92	0.90	1000
9	0.88	0.89	0.89	1000
micro avg	0.82	0.82	0.82	10000
macro avg	0.82	0.82	0.82	10000
weighted avg	0.82	0.82	0.82	10000

Training Accuracy vs Validation Accuracy



Training Loss vs Validation Loss



MODEL COMPARISON



Traditional Models

	Random Forest	Logistic Regression	Gradient boost
Without PCA	R ² Train Data: 0.64 R ² Test Data: 0.43	R ² Train Data: 0.52 R ² Test Data: 0.36	R ² Train Data: 0.94 R ² Test Data: 0.44
Time	00:08:16.12	01:49:33.14	09:24:40.93
With PCA	R ² Train Data: 0.60 R ² Test Data: 0.43	R ² Train Data: 0.38 R ² Test Data: 0.38	R ² Train Data: 0.92 R ² Test Data: 0.44
Time	00:01:53.35	00:00:16.38	00:09:43.77

Convolutional Neural Networks

	Model1	Model1b	Model2b	Model3	Model4
Loss	1.72	1.25	0.89	0.886	0.830
Accuracy	0.38	0.55	0.68	0.814	0.816
Time	02:29:19.31	02:18:07.45	00:09:43.77	05:28:11.04	10:30:20.44

- Results shown were obtained after training and testing the models at least 3 to four times with different parameters.
- Using Principle Component Analysis (PCA) for dimensionality reduction appears to provide time efficient results without compromising accuracy. Same results that were achieved using regular modeling techniques and in some cases took hours to process can be achieved in fraction of time using PCA.
- Likewise for convolutional neural networks as the network continues to grow it puts a heavy penalty on processing times as more layers are added to it. Therefore we require more computational power for tuning different parameters and achieving desired results in specified time.
- In conclusion the models that serve the purpose well for image recognition include Convolutional Neural Networks and Random Forest. CNN although time consuming offers far better accuracy than any traditional machine learning model.
- For this project CNN (Model3) provides best possible results with a loss of 0.886 and 81% accuracy with a processing time of 5 hours and 30 min followed by another CNN (Model4) which takes in more layers and nonetheless provides same level of accuracy within twice the amount of time which is 10 hours 30 min.